


Nov 22, 19 18:35

Proj2_Spec.txt

Page 1/11

```
*****
***                      COMP 15 Project 2                      ***
*****
```



May you find what you are searching for
in the right places
âM-^@M-^U Lailah Gifty Akita

```
//////////////////////////////////////
//                                     //
//          Gerp - It's Like grep But Something Is Off          //
//                                     //
//////////////////////////////////////
```

We're all familiar with web search engines, and we also have tools for searching our personal computers. Have you ever wondered how the Mac Spotlight works, for example? We'll look at one approach now!

In this assignment you will implement a program that indexes and searches a file tree for strings. Your program will behave similarly to the Unix "grep" program, which can be used like this to search through all the files in a directory and look for some sequence of characters:

```
grep -Rn Query DirectoryToSearch
```

Where "Query" is the target string, and "DirectoryToSearch" is the directory we want to look in for the "Query".

For example,

```
grep -IRn ifstream /comp/15/files
```

Will produce:

```
/comp/15/files/lab1/main.cpp:45:      ifstream in;
/comp/15/files/extras/read_from_file.cpp:34: *      cin or an ifstream.
/comp/15/files/extras/read_from_file.cpp:40: *      of ifstream a C++ string.
/comp/15/files/extras/read_from_file.cpp:56:      ifstream input;
/comp/15/files/lab4/roster.cpp:51:      ifstream infile;
/comp/15/files/lab7/er.cpp:45:      ifstream infile;
```

This output tells you that "ifstream" occurs on line 45 /comp/15/files/lab1/main.cpp and on lines 34, 40 and 56 of /comp/15/files/hw3/read_from_file.cpp, for example, and it prints out the matching line after the colon.

How can we do something like this?

Your program will 1) build an index data structure and 2) use that index to respond to queries.

Nov 22, 19 18:35

Proj2_Spec.txt

Page 2/11

When you index files, you create a database containing information on all the files within the directory. Queries are searches that a user want you to do, e. g., "tell me all files that have the word "potato."

To create an index, you will read in the files and store information about them (such as their names, their relative paths, and their contents) in a data structure of your choice that is easily searched and/or queried.

We'll first describe the program, then the implementation details, and then the submission details.

Please read through the entire assignment before you start.

Program specification

Your program will traverse a file tree created using a module that we provide (see the interface and description of the file tree in "Using the FSTree and DirNodes" under "Implementation Details" below). It will need to index each file that it finds in the tree. After indexing all of the files it will enter a command loop (similar to the "interactive" modes that you have implemented in previous HWs) where the user can enter various commands to modify the search, and to quit the program.

Your program will be started from the command line like this:

```
./gerp DirectoryToIndex OutputFile
```

"DirectoryToIndex" determines which directory will be traversed and indexed. You could use "/comp/15/files", for example. "OutputFile" names the file to which query results should be sent.

Once the program indexes the specified directory it will print "Query?" on cout (NOT to the OutputFile) and wait for a command from the user. The possible query commands are the following:

- o AnyString

A word (see "What is a Word?" under "Implementation Details") is treated as a query. The program will take this string and print all of the lines in the indexed files where "AnyString" appears. Note this a case sensitive search so "we" and "We" are treated as different strings/words and should have different results.
- o "@i AnyString" or "@insensitive AnyString"

Preceding a query string by "@i" or "@insensitive" causes the program to perform a case insensitive search on the string that was passed. For example, "we" and "We" would be treated as the same string/word and will have the same results.
- o "@q" or "@quit":

These commands will completely quit the program, and print "Goodbye! Thank you and have a nice day." This statement should be followed by a new line. Note that the program should also quit when it reaches End-Of-File (EOF).
- o "@f newOutputFilename"

This command causes the program to close the current output file. Any future output should be written to the file named newOutputfilename.

NOTE: Commands will be entered without the quotation marks ("")

If the user did not specify exactly two command line arguments (in addition to the program name), print this message on cerr:

Nov 22, 19 18:35

Proj2_Spec.txt

Page 3/11

Usage: gerp directory output_file

and terminate the program (by returning EXIT_FAILURE from main() or by calling exit(EXIT_FAILURE)).

BEWARE: There are two output streams.

- o The "Query? " prompt always goes to cout.
- o The results of the query always go to the output file, which'll either be OutputFile or the file name in the last "@f" command.

Treat a multi-word query as several independent 1-word queries.
E. g.

Query? We are the champions

is the same as

Query? We
Query? are
Query? the
Query? champions

For this assignment, Standard Template Library (STL) implementations may be used. A description of STL implementations that you are allowed to use are listed in their own section "STL Usage" under "Implementation Details".

To help you learn the interface and get a feel for the program, we have provided you with a working reference implementation.

The Files to Implement

For this assignment we will not specify the files or functions you will need to write. Instead your program is required to function as described in this specification. You may accomplish this task using any combinations of files, functions, and classes you wish. We will, of course, evaluate your design.

In addition to writing .h and .cpp files for your classes you will need to write a main() for your program, and write a Makefile. The default make action should be to compile and link the entire program producing an executable program named "gerp", which you can run by typing "./gerp"

If you have a "clean" target in your Makefile, be sure you do not delete the .o files we give you to.

You will want to write test methods to test the various parts of your program separately so that you do not have to debug compounded errors.

Nov 22, 19 18:35

Proj2_Spec.txt

Page 4/11

Implementation Details

What is a Word?

It is important to outline what a word is when dealing with a word search engine. We will define a word as a string that starts and ends with an alphanumeric (letter or number) character. This means that you will need to do a little string parsing to determine the output of your gerp implementation. To help you with this nuance we have included a couple examples.

When searching for the word "comp" using case insensitive sort, gerp should treat the following strings as the same as

```
"comp":
  o "comp"
  o "comp."
  o "Comp"
  o "-comp"
  o "&&comp"
  o "comp?!"
  o "@#$comp?!"
```

Also if any of the bulleted strings were submitted as a query, gerp should print lines in files that contain any of the strings on the list (however it should print them as they exist in the file, not a processed version).

NOTE: gerp should compare strings where all leading and trailing non-alphanumeric characters are striped. This includes both the queries and the the strings in the data files.

Output Formatting

If the word in a query is found in the index, then, for each line it appears in, you will print to the designated output file a line of the form

```
FileNameWithPath:LineNum: Line
```

which consists of:

1. The full pathname of the file (including the path from the command line), followed by a colon
2. The line number within that file the query word appears on, followed by a colon and a space
3. The full text of the line from the file
4. A newline

For example if you ran

```
$ ./gerp small_test out.txt
Query? we
```

which queries "we" on our small_test directory, and then opened the file out.txt, you might see

```
small_test/test.txt:5: we are the champions
small_test/test.txt:6: we we we
```

NOTE: THERE IS ONE NEW LINE AFTER THE LAST LINE.
EACH LINE THAT THE QUERY APPEARS ON ONLY PRINTS ONCE.

If the query is not found using the default search your implementation should print:

```
query Not Found. Try with @insensitive or @i.
```

Nov 22, 19 18:35

Proj2_Spec.txt

Page 5/11

If the query is not found using the insensitive search your implementation should print:

query Not Found.

NOTE: there is no whitespace before the messages, which begin in column 1.

Using FSTree and DirNodes

FSTree: We use a file-system tree to represent directories, subdirectories, and files. The data structure we use is an n-ary tree, so-called because a node of the tree could have any number of children. The main usage of this class is to help you navigate through folders and directories inside a computer.

For example, a snapshot of a home directory might be represented in a n-ary tree like this:

```
/h/mkorman
  /comp11
    /test/comp15
      /exams /lab
      /assignments
      /hws /projects /comp160
      /hw
```

YOU DO NOT HAVE TO WRITE THIS CLASS --- We did it for you.

An FSTree is an n-ary tree consisting of DirNodes (which will be described below). The FSTree class has the following public functions:

- o FSTree(string rootName)
This is the constructor for an FSTree. It creates a file tree of DirNodes where the root of the tree is the directory that was passed as a parameter. If there is an error opening directories or files, the constructor will fail and halt your program. Be careful and do not run this on just any directory: if a directory structure has a loop in it, the constructor can run forever.
- o ~FSTree()
The destructor deallocates all of the space allocated when the tree was built.
- o void burnTree()
This function destroys the tree, frees any data that was allocated, and makes the tree empty.
- o DirNode *getRoot()
This function returns the root of the tree. Normally we do not want to return the private members of an object or class, however in this case it is necessary so that you can traverse the tree and index its contents.

The DirNode class is a building block for the FSTree class. It is our representation of folders. Each DirNode instance has a string name, list of files in the directory, and a list of subdirectories. It contains the following public methods:

- o bool hasSubDir()
returns true if there are sub directories in the current node (directory).
- o bool hasFiles()
returns true if there are files in the current node (directory).
- o bool isEmpty()
returns true if there are no files or sub directories

Nov 22, 19 18:35

Proj2_Spec.txt

Page 6/11

- in the current node (directory).
- o int numSubDir()
returns the number of sub directories.
- o int numFiles()
returns the number of files in the current node.
- o string getName()
returns the name of the current directory.
- o DirNode *getSubDir(int n)
returns a pointer to the nth subdirectory.
- o string getFile(int n)
returns nth file name
- o DirNode *getParent()
get parent directory node

The DirNode class also contains the following public functions which are used to modify the contents and structure of the tree (we use these functions to initially build the FSTree). Although you have access to these functions, you should refrain from using them, as it may cause the tree to lose data/information. We're only telling you about them, because you'll see them listed in the .h file, and you may be curious.

- o DirNode(string newName)
Constructor that initializes a DirNode named newName.
- o void setName(std::string newName)
set the name of the current node.
- o void addFile(std::string newName)
Adds a file with the name "newName" to the current node.
- o void addSubDirectory(DirNode *newDir)
Adds a sub directory (newDir) to the current node.
- o void setParent(DirNode *newParent)
Sets parent node (directory) of the current node.

In order to get a file's full path you will need to traverse the FSTree and concatenate the names of the directories you traverse to compile a file's full path (which is necessary to subsequently index the file). You will then use this full path to open the file in an ifstream and index its contents.

STL Usage

For this assignment you will be allowed to use *ONLY* the following STL template implementations:

- o vector
- o queue
- o stack
- o set
- o unordered_set
- o functional

You are not required to use any particular items in the STL. If you feel one would be useful, you will need to learn more about their respective interfaces. You can find more information at:

<http://www.cplusplus.com/reference/>

Any other data structures you need, you must implement yourself.

Compiler Optimizations

When compiling your implementation of "gerp" you should compile it with the flag "-O2". (That's a capital letter 'O', not the numeral zero.) This will optimize your program for the system that it is compiling on, which will result in an implementation with a faster run time. This will help everyone during the testing phase because one will receive results faster.

Nov 22, 19 18:35

Proj2_Spec.txt

Page 7/11

```
# Testing and Reference Implementation#
In order to help you with your testing and to get familiar with the
user interface expectations of this project. We have provided you
with a fully compiled reference implementation called "the_gerp".
By the end of the project your gerp implementation should behave
exactly the same as the_gerp.
```

```
*****
* NOTE: YOUR VERSION OF GERP SHOULD BEHAVE EXACTLY AS THE *
* REFERENCE MATERIAL WE SUPPLIED ("THE_GERP") UNDER ALL *
* CIRCUMSTANCES. THEREFORE YOU SHOULD EXTENSIVELY TEST THE *
* REFERENCE TO FIGURE OUT HOW IT BEHAVES. *
*****
```

You can easily compare your output to the reference by redirecting output to two different files:

```
./the_gerp Directory ref_output.txt < commands.txt
```

This command redirects the program's input so that the contents of "commands.txt" appear on cin just as if someone typed it in.

You could run your program similarly:

```
./gerp Directory my_output.txt < commands.txt
```

Then sort both text files using the unix "sort" command:

```
sort ref_output.txt > ref_output_sorted.txt
sort my_output.txt > my_output_sorted.txt
```

After sorting you can use the "diff" command to find the differences between the two files:

```
diff ref_output_sorted.txt my_output_sorted.txt
```

"diff" will print the differences (if there are any) to the terminal. If nothing prints out, the files are identical.

Sorting is necessary, because, while your program must produce output for all occurrences of the query, the order in which multiple lines appear is not specified. Print out multiple lines in whatever order your data structure and algorithm choices find convenient.

Building/Indexing and Queries

When designing and implementing your program you should aim to have it build its index and run queries as quickly as possible. You may find that there is a trade off between the two (e.g. a program that builds an index quickly may not search as fast). It is important to document your design choices, your justification of the choices, and their effects in your README.

To get full credit, your program will need to be able to index the largest file tree and be ready to query in under 10 minutes. Our implementation uses approximately 2.3 GB of RAM, once it has fully indexed the file tree. On our system, your program can use about 7.5 -- 8 GB of memory before it crashes or is killed by the operating system. This means that you will lose points for tests on the largest collection if you exceed the memory requirements. You can get credit for tests on the smaller collections, of course, if your program works for those.

You can check the memory usage of your program using the following unix command:

Nov 22, 19 18:35

Proj2_Spec.txt

Page 8/11

```
echo "@q" | /usr/bin/time -v ./gerp [DirectoryToIndex]
```

This command will start your program and it will complete its indexing procedure then immediately quit. In the output, "Elapsed (wall clock) time (h:mm:ss or m:ss)" is the duration that it took to index the directory. "Maximum resident set size (kbytes)" is the peak memory usage of your program. Do not forget that the results are in KB and will need to be converted to GB (multiply by 1,000). If you want to test our program's indexing speed you should replace "./grep" with "./the_gerp".

Also please note our reference implementation, "the_gerp", is by no means the fastest indexing or fastest querying solution to the problem. We hope that you are able to build a faster implementation!

Other Implementation Details

DO NOT IMPLEMENT EVERYTHING AT ONCE!

This may seem like a lot, but if you break it into pieces, it's perfectly manageable. Just do it one bit at a time.

Each class you design should do one thing. Clear abstractions with clearly delineated responsibilities are easier to design, implement, test, understand, and debug. Therefore, it is as important for you as for us --- you'll work more efficiently and have fewer bugs!

You will add one function, then write code in your test file that performs one or more tests for that function. Write a function, test a function, write a function, test function, ... This is called "unit testing."

Follow the same testing approach for every class you write!

You should add functionality little by little.

```
*****
* NOTE: YOU WILL NOT BE SUBMITTING A TESTING MAIN ALONG WITH *
* YOUR GERP IMPLEMENTATION. HOWEVER YOU ARE REQUIRED *
* TO DETAIL YOUR TESTING METHODS IN YOUR README *
*****
```

If you need help, TAs will ask about your testing plan and ask to see what tests you have written. They will likely ask you to comment out the most recent (failing) tests and ask you to demonstrate your previous tests.

Be sure your files have header comments, and that those header comments include your name, the assignment, the date, the purpose of the particular file, and acknowledgements for any help you received.

```
-----
Makefile
-----
```

In addition to the other required files you will need to submit a Makefile that compiles your program when the user types "make" in the terminal. The name of your fully compiled program should be "gerp".

```
-----
README
-----
```

With your code files you will also submit a README file. You can format

Nov 22, 19 18:35

Proj2_Spec.txt

Page 9/11

your README however you like. However it should have the following sections:

- A. The title of the homework and the author's name (you)
- B. The purpose of the program
- C. Acknowledgements for any help you received, including references to outside sources you consulted (though there is no need list C++ references like cplusplus.com).
- D. The files that you provided and a short description of what each file is and its purpose
- E. How to compile and run your program
- F. An "architectural overview," i. e., a description of how your various program modules relate. For example, the FSTree implementation keeps a pointer to the the root DirNode.
- G. An outline of the data structures and algorithms that you used. Given that this is a data structures class, you need to always discuss the the data structure that you used and justify why you used it. For this assignment it is imperative that you explain your data structures and algorithms in detail, as it will help us understand your code since there is no single right way of completing this assignment.
- H. Details and an explanation of how you tested the various parts of your classes and the program as a whole. You may reference the testing files that you submitted to aid in your explanation.

Each of the sections should be clearly delineated and begin with a section heading which describes the content of the section.

Submitting Your Work

You will be submitting your work in 2 parts as described below:

- o Part1 - Design + Using the FSTree and String Manipulation:
 - Design

Because HW 5 in-person grading is on-going, we will not do in-person check-offs for your Project 2 design. Rather, you will turn a design.text file in with provide that addresses these items:

 - * What data structure(s) will you be using? Why these data structures?
CAUTION: If your answer is just vectors, you need to think again! Use the data structures we've been studying.
 - * What classes will you implement (presumably including one for each data structure above)? What public functions does each class support?
 - * How do your classes interact, i. e., does one class contain an instance of another, a pointer to an instance of another, what functions in the other class will it call?
 - * How will you implement insert, search?
 - * Describe the space needs of your solution and the big-O runtime of important operations (insert, search).
 - * You may optionally submit an image file with a diagram of your implementation, but keep it under 1 MB in size.
 - FSTree and String Manipulation

To complete this part you will need to write 2 programs, each with its own main function: (1) a tree traversal method that prints out the full paths of each file in the tree on separate lines, (2) a function that strips all leading and trailing non-alphanumeric characters from a given string.

Nov 22, 19 18:35

Proj2_Spec.txt

Page 10/11

You should write your tree traversal in a file called "FSTreeTraversal.cpp". This program should take the highest directory as a command line argument (see below) and then print the full paths of all of the files accessible from that directory:

```
./treeTraversal Directory
```

Do not worry about the order that the file paths print in, just ensure that each one of them prints.

Your string processing method should be defined in a file called "stringProcessing.cpp" and have a declaration in a file called "stringProcessing.h". The function should be called "stripNonAlphaNum", it takes a string as a parameter and returns a string. This function should remove all leading and trailing non-alphanumeric characters such that when "@###%#COMP-15!!!!!!" is given as a parameter the function returns "COMP-15".

In addition to the defining the stripNonAlphaNum function, you will want a main function in stringProcessing.cpp for testing. Write complete a program that reads strings from cin (until reaching end of file), calls stripNonAlphaNum on each string and prints the result on cout, each result string on a line by itself. Remember: stripNonAlphaNum should be a separate function.

The provide command is:

```
provide comp15 proj2part1 FSTreeTraversal.cpp \
stringProcessing.h \
stringProcessing.cpp \
README design.text
```

You may include an image file with a diagram of your design (under 1 MB in size). Your README does not have to contain everything we normally ask. For this phase, you can just write a very brief summary with any information you want the grader to see, e. g., if you have a nettlesome bug you have not been able to fix.

- o Part2 - Final Submission:

For this part you will submit all of the files required (including a Makefile) to compile your gerp program. The filenames for any testing modules you chose to submit should begin with "test". The provide template is:

```
provide comp15 proj2part2 README Makefile [YOUR FILE NAMES HERE]
```

NOTE: We cannot give you a complete provide command --- so make sure you submit everything we will need to run your program. Maybe copy the files into another directory, type "make" test the program, do a "make clean" and then provide every thing. We should be able to use "make" or "make gerp" to build your program.

Nov 22, 19 18:35

Proj2_Spec.txt

Page 11/11

----- Helpful tricks -----

The data sets we're using include a very large collection. Do not copy them. You can pass them directly to your program:

```
./gerp /comp/15/files/proj2-test-dirs/largeGutenberg output.txt
```

or you can link them into your directory like this:

```
ln -s /comp/15/files/proj2-test-dirs ./test-dirs
```

which will create a "link" in your directory that refers to our directory. You can use it as if it were a local directory without copying over gigabytes of data.

We will test your solution on several directories, including smallGutenberg, mediumGutenberg, and largeGutenberg. You can see some test queries and reference implementation output in /comp/15/files/proj2-sample-execution. You should match this output exactly (except that the order can be different as explained above).

----- Data structures -----

You are welcome to use data structures of your choosing, those we've covered in class and those we have not. You are allowed to use the STL data structures listed under "# STL usage #". All other data structures must be implemented yourself. Do not forget to cite any source that you use in your acknowledgements.

There are example uses of the hash facility in the files we've given you. See hash_test.cpp and hashExample.cpp.

If you use a hash table, then you MUST have it dynamically resize. I. e., you must monitor the load factor and expand the hash table if the load factor is exceeded.

The runtime of your query processing must depend on the size of the output, not on the size of the input (that is, the index structure should be essentially constant time to search).

Finally, space will be an issue. Your program will have enough memory to store all the data in the largeGutenberg collection once or twice, but not 5 or 6 times. If you are not cognizant of space usage, then your program will work fine for the small or maybe even the medium collection, but will fail on the large collection when it runs out of memory.

If you are stuck, then you should get something that at least works correctly on the small and medium collections. But to get a very high score, it will also need to work on the larger collection, and that means choosing a strategy that considers space.

Hint: Every copy of a string requires space proportional to the length of the string. (A string is an ArrayList of characters.)

Hint: Do not store a pointer to an element in a vector. A good final exam question would be: "Why should you not store the address of an element in a vector?" A vector is an ArrayList, remember.